

MAHAD REHMAN DURRANI

Islamabad, Pakistan

+92 333 9131317 | mahadrehman04@gmail.com | [in mahad-rehman](https://www.linkedin.com/in/mahad-rehman) | [mahadrehmann](https://github.com/mahadrehmann)

EDUCATION

Bachelor of Science in Computer Science
FAST National University (NUCES), Islamabad

Expected June 2026

EXPERIENCE

AI Engineer (Apprentice)

XFlow Research — June 2025 - April 2026

- Engineered production-grade **Agentic Workflows** to automate enterprise processes, focusing on scalability.
- Reduced LLM hallucinations and inference costs by implementing strict **RAG evaluation loops**.
- Gathered and translated client requirements into actionable technical specifications.

Teaching Assistant (Generative AI & Software for Mobile Devices)

FAST NUCES

- Conducted hands-on technical demos and mentored students in implementing advanced **Generative AI architectures**, including Transformers, GANs, VAEs, Diffusion models, etc.
- Reviewed architectures and API integration strategies for 50+ **Mobile Applications** (Kotlin/Flutter).

PROJECTS

Markly: Autonomous Marketing Agent System

FYP — FastAPI, LangGraph, Docker

- Engineered a multi-agent marketing system automating the end-to-end campaign lifecycle from **market research** to **social media posting**.
- Built a Computer Vision & GenAI pipeline (YOLOv8, CLIP, Flux.1, Supabase) for data-driven competitor analysis and automated poster generation.

RAG Chatbot with Multi-Stage Evaluation Pipeline

GenAI — LangChain, OpenAI, Gemini

- Built a ReAct-style RAG agent with **LangChain** and FAISS, fortified by a strict input validation layer using the OpenAI Moderation API and Gemini.
- Engineered a 4-stage **"LLM-as-a-Judge"** pipeline using LangSmith to trace executions, evaluate factual consistency, and mathematically track pass rates.

MailGen: Agentic Email Automation System

Agentic AI — CrewAI, Flask

- Orchestrated a multi-agent workflow using **CrewAI** to autonomously analyze requirements and generate context-aware professional emails.
- Integrated a **Flask** web dashboard with background job scheduling to handle automated bulk dispatch and dynamic file attachments.

End-to-End MLOps Pipeline

AWS, GitHub Actions, Prometheus, Docker

- Engineered a production-grade MLOps system with **automated model retraining**, **real-time observability** (Prometheus/Grafana) with Slack alerting, and containerized CI/CD deployments to AWS EC2 via GitHub Actions.

CERTIFICATIONS

- **Machine Learning Specialization** by DeepLearning.AI
- **AWS Academy Cloud Foundations** by AWS
- **Generative AI with LLMs** by DeepLearning.AI

TECHNICAL SKILLS

Languages	Python, C/C++, SQL, Java, Bash
AI/ML	Machine Learning, Deep Learning, Generative AI, RAG Pipelines, Agentic Systems
Frameworks & Libs	PyTorch, TensorFlow, HuggingFace, LangChain, LangGraph, CrewAI, FastAPI, NumPy, Pandas, Scikit-Learn, Matplotlib
DevOps & Cloud	AWS, GCP, Docker, Kubernetes, GitHub Actions, Jenkins, Prometheus, Grafana, Linux/WSL